

## Question 1: Explain the fundamental differences between DDL, DML, and DQL commands in SQL.

Provide one example for each type of command.

**Answer:**

The fundamental difference between DDL, DML, and DQL commands lies in their **purpose and scope** within the database system.

Command Type	Full Form	Purpose	Example
DDL	Data Definition Language	Used to <b>define</b> or modify the structure (schema) of database objects (tables, indexes, views, etc.). These commands affect the database structure itself.	<code>CREATE TABLE table_name (column_name data_type);</code>
DML	Data Manipulation Language	Used to <b>manage</b> the data within the schema objects, allowing users to insert, update, or delete records. These commands affect the data stored in the tables.	<code>INSERT INTO table_name (column1) VALUES (value1);</code>
DQL	Data Query Language	Used solely to <b>retrieve</b> data from the database. It consists only of the <code>SELECT</code> command.	<code>SELECT column1, column2 FROM table_name WHERE condition;</code>

---

## Question 2: What is the purpose of SQL constraints?

Name and describe three common types of constraints, providing a simple scenario where each would be useful.

**Answer:**

The purpose of **SQL Constraints** is to enforce rules on the data columns of a table to limit the type of data that can go into a table. They ensure the **accuracy** and **reliability** of the data in the database.

### Three Common Types of Constraints:

1. **PRIMARY KEY**
  - **Description:** A field or set of fields that uniquely identifies each record in a table. It must contain unique values and cannot contain NULL values.
  - **Scenario:** In a `Students` table, the `StudentID` column should be the PRIMARY KEY to ensure that every student has a unique identifier and no record is duplicated.
2. **FOREIGN KEY**
  - **Description:** A field in one table that uniquely identifies a row of another table, establishing a link between the two. It enforces **referential integrity**.
  - **Scenario:** In an `Orders` table, the `CustomerID` column should be a FOREIGN KEY referencing the `Customers` table. This ensures that an order can only be placed by an existing customer.
3. **NOT NULL**
  - **Description:** Ensures that a column cannot have a **NULL** value, guaranteeing that data is always present in that column for every record.
  - **Scenario:** In a `Employees` table, the `LastName` column should be NOT NULL because it is essential information for identifying an employee.

---

### Question 3: Explain the difference between LIMIT and OFFSET clauses in SQL.

How would you use them together to retrieve the third page of results, assuming each page has 10 records?

**Answer:**

- **LIMIT Clause:** Used to **restrict** the number of rows returned by a query, specifying the maximum number of rows to retrieve.
- **OFFSET Clause:** Used to **skip** a specified number of rows from the beginning of the result set before starting to return the required result rows.

#### Retrieving the Third Page of Results

To retrieve the third page of results with 10 records per page, we need to skip the first two pages ( $2 \times 10 = 20$  records).

- **LIMIT** (Page Size) = 10
- **OFFSET** (Records to Skip) = 20

The query is:

SQL

SELECT columns

FROM table\_name

ORDER BY some\_column -- ORDER BY is crucial for consistent pagination

```
LIMIT 10      -- Retrieve 10 records (the third page)
OFFSET 20;    -- Skip the first 20 records (pages 1 and 2)
```

---

## Question 4: What is a Common Table Expression (CTE) in SQL, and what are its main benefits?

Provide a simple SQL example demonstrating its usage.

**Answer:**

A **Common Table Expression (CTE)** is a temporary, named result set that you can reference within a single SQL statement (e.g., `SELECT`, `INSERT`, `UPDATE`, or `DELETE`). It is defined using the `WITH` clause.

### Main Benefits:

1. **Readability:** CTEs break down complex queries into simpler, logical, and more readable blocks.
2. **Reusability:** The CTE can be referenced multiple times within the same query.
3. **Recursion:** CTEs are necessary for performing recursive queries (queries that reference themselves).

### Simple SQL Example:

This example calculates the average price of all products (in the CTE) and then selects products that are above that average.

SQL

```
WITH AveragePrice AS (
    -- Define the CTE: Calculate the average price
    SELECT AVG(Price) AS AvgProductPrice
    FROM Products
)
-- Use the CTE: Select products above the calculated average
SELECT ProductName, Price
FROM Products
WHERE Price > (SELECT AvgProductPrice FROM AveragePrice);
```

---

## Question 5: Describe the concept of SQL Normalization and its primary goals.

Briefly explain the first three normal forms (1NF, 2NF, 3NF).

**Answer:**

**SQL Normalization** is the systematic process of organizing the columns and tables of a relational database to minimize **data redundancy** and improve **data integrity**.

### **Primary Goals:**

1. **Eliminate Redundant Data:** Avoid storing the same data in multiple places to save storage and prevent inconsistencies.
2. **Ensure Data Dependencies Make Sense:** Ensure that the data is logically stored, meaning that only related data is grouped together in a table.

### **The First Three Normal Forms:**

1. **First Normal Form (1NF):**
  - The table must contain only **atomic** (indivisible) values in each column.
  - There must be no repeating groups of data.
2. **Second Normal Form (2NF):**
  - The table must be in **1NF**.
  - There must be **no partial dependencies**. This means no non-key attribute can depend on only a *part* of a composite primary key.
3. **Third Normal Form (3NF):**
  - The table must be in **2NF**.
  - There must be **no transitive dependencies**. This means non-key attributes must not depend on other non-key attributes; all non-key attributes must depend directly on the primary key.